

Introduction to Tensorflow in Scala

Xavier Tordoir

Schiphol Developers Group


13 Feb 2020



LUNATECH

SIMPLIFY YOUR IT

Introduction to **Tensorflow in Scala**



Xavier Tordoir, Lunatech



LUNATECH
SIMPLIFY YOUR IT



JVM development

Devops

ML & Big Data

 **Scala**

 **TensorFlow**

 **play**

 **Spark**

 **kafka**

 **amazon
web services™**

 **DC/OS**

 **Java**

 **akka**

 **kubernetes**

 **ActiveMQ**



 **Kotlin**

 **cassandra**

 **LUNATECH**

105

Employees

27

Nationalities

Lots

Open source

Rotterdam NL



We welcome
you to one
of our offices

Paris FR



Amsterdam NL



 **LUNATECH**

Outline

ML Concepts

Scala notebooks: Interactive Programming


Tensorflow Framework, Overview with Scala Tensorflow

Inference service with Akka http

Managing ML artefacts

Managing development flow

ML Introduction



ML Introduction

Data as a flat table

```
type Feature = Double
type Label   = Double

val dataSet: Seq[ (Vector[Feature], Label) ]
```

Surface	Land	Beds	Sidings	Price
110	896	2	4	160
120	435	3	2	189
150	210	4	3	250
170	713	4	4	240
80	231	4	4	179
90	238	3	4	135
130	118	2	3	175
146	695	4	4	169
155	644	4	4	189

ML Introduction

A model is function a representing a facet of the data

```
val model: Vector[Feature] => Label
```

Surface	Land	Beds	Sidings	
110	896	2	4	

 \Rightarrow

Price
180

ML Introduction

Learning a Model from Data

```
val train: Seq[ (Vector[Feature], Label)] => Vector[Feature] => Label
```

Surface	Land	Beds	Sidings	Price
110	896	2	4	160
120	435	3	2	189
150	210	4	3	250
170	713	4	4	240
80	231	4	4	179
90	238	3	4	135
130	118	2	3	175
146	695	4	4	169
155	644	4	4	189

=>

Surface	Land	Beds	Sidings
110	896	2	4

=>

Price
160

ML Introduction

Training by Minimizing Errors (Loss), e.g. sum of squared errors:

```
val loss = dataSet.map{  
  case (x, y) => y - model(x)  
}  
  .map(Math.pow(_, 2))  
  .reduce( _ + _ )
```

Surface	Land	Beds	Sidings
110	896	2	4
120	435	3	2
150	210	4	3
170	713	4	4
80	231	4	4
90	238	3	4
130	118	2	3
148	695	4	4
155	644	4	4

Price
160
189
250
240
179
135
175
169
189

Price
160
189
250
240
179
135
175
169
189

$$\text{Loss} = \sum_i (y_i - \hat{y}_i)^2$$

Missing pieces yet: How a model is built? What is 'minimizing'?

ML Introduction

Models as a vector of parameters


A model is a function, with some parameters, optimisation is finding the best parameters...

Example: A **Linear model** is a linear combination of features:

```
val params: Vector[Double]
val bias: Double

val model = x: Vector[Feature] =>
  x.zip(params).map( case(wi, xi) => wi * xi )
    .reduce( _ + _ )
    + bias
```

$$\hat{y} = \sum_i x_i w_i$$



ML Introduction

Optimisation algorithms

Gradient based methods: How loss varies with each parameters ~ gradient ()

$$\Delta Loss \sim \Delta \omega_i$$

$$\omega_i^* = \omega_i - \gamma \frac{\Delta Loss}{\Delta \omega_i}$$

Loss and gradient are estimated on a subset of data (a batch) = stochastic gradient based methods

Iterations in batches and epochs (a full dataset pass)

ML Introduction

Metrics

After training: model evaluation

E.g.

Root Mean Squared Error in regression

Accuracy in classification (% correct binary prediction)

Metrics are used for model validation on test data not used in training

Surface	Land	Beds	Sidings	Price
110	896	2	4	160
120	435	3	2	189
150	210	4	3	250
170	713	4	4	240
80	231	4	4	179
90	238	3	4	135
130	118	2	3	175
146	695	4	4	169
155	644	4	4	189

ML Introduction

Data is multidimensional **Arrays** of **Floating** point values

Models are represented as **Arrays** of **Floating** point values and operators

Training, **Evaluating** and **Inference** on models are **operations** on these arrays

Tensorflow: ML Framework

Computing paradigm for **Tensors** transformations

ML capabilities, Linear models and Neural networks

Tensorflow: tensorflow_scala

Comprehensive tensorflow library in Scala:

https://github.com/eaplatanios/tensorflow_scala

http://platanios.org/tensorflow_scala/

<https://brunk.io/deep-learning-in-scala-part-3-object-detection.html>

JNI layer to interface the c library

Tensorflow Framework generated with protobuf

Scala API mimics lots of the python API

Reference scala API

Tensorflow: Notebooks

Data Science implies:

- **knowledge** of the data, including its corner cases
- **Exploration** of how to guide the modelling, choosing the right methods for the data
- **Trials and errors**, no possibility to implement functional specs, only model validation

=> Need for **interactive programming**

=> **Notebooks**

Tensorflow: Scala Notebooks



<https://jupyter.org/>

- Python environment
- Kernels to support different languages:

<https://almond.sh/>



Zeppelin

<https://zeppelin.apache.org/>

Spark-notebook

<http://spark-notebook.io/>

Scala notebooks

Tensorflow: Jupyter Notebooks Environments

Environments

Local install (Linux with GPU or OSX on CPU)

- Virtualenv for python
- Jupyter
- coursier
- almond to wrap Ammonite

Google colaboratory (CPU, GPU, TPU)

- Preinstalled python/jupyter
- Console for Scala kernel

Custom Server & other providers

Tensorflow: Overview of the core

Tensor as a Multidimensional array

Variables and Placeholders

Operators

Session and Graph

Abstraction from computing device and Linear Algebra Accelerator

Tensorboard

Linear Regression with Low Level API

Tensorflow Scala Examples

Very simple Linear Regression example

00_tf_intro.ipynb

Tensorflow: Tensors

Tensor \Leftrightarrow Multidimensional array

Rank = #dimensions

Shape = sizes for each dimension

Type = Type of data (Int, Double, String)

Static Tensors manipulation, read a csv to tensor

Tensorflow: Dataset API

Unified API over multiple Data sources

Nested structure of Tensors (like a collection)

Iterators => basis for batching

Dataset API provides encoding and decoding of Text, Images, Protobuf (TFRecords)

Operations like batches, repeats, shuffle

Sharding for distributed training

Load a Tensor as a Dataset

Tensorflow: Estimators API

Encapsulates everything needed for training a Model:

Model Signature:

- Input to define input types (like a placeholder shape)
- TrainInput to define prediction types (Again like a placeholder)

Model Function and Parameters:

- Build Neural network topologies with a simple API
- Trainable Parameters

Loss function, Optimizer, Metrics

Hooks for:

- Save models regularly (Graph structure and parameters) — checkpoint files
- Log metrics and Tensorboard

Define Model and Training

Tensorflow Scala Examples

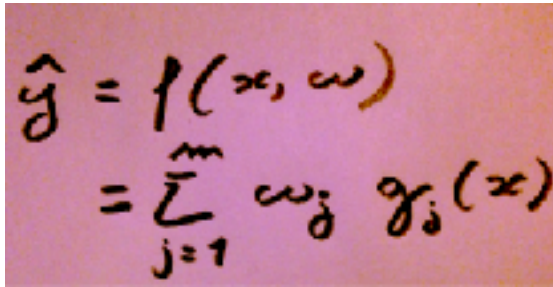


Using Datasets and Estimators

01_estimators_datasets.ipynb

Linear Models to Deep Learning


Linear Model


$$\hat{y} = f(x, w) = \sum_{j=1}^m w_j g_j(x)$$

```
Linear[Float]("Layer_0", 1)
```

Multi Layer Perceptron

```
Linear[Float]("Layer_0", 128)>> ReLU[Float]("Layer_0/Activation", 0.1f) >>  
Linear[Float]("Layer_1", 64) >> ReLU[Float]("Layer_1/Activation", 0.1f) >>  
Linear[Float]("Layer_2", 32) >> ReLU[Float]("Layer_2/Activation", 0.1f) >>  
Linear[Float]("OutputLayer", 10)
```




Deep is defined by accumulation of layers...and parameters

Convolutional Neural Networks

```
tf.learn.Conv2D("Layer_1/Conv2D", Shape(3, 3, 3, 32), stride1 = 1, stride2 = 1, padding = ValidConvPadding) >>
```

Convolution Layer




**Each Pattern is scanned against the image and
scores are computed in the layer
Training learns the patterns that work well...**

Convolutional Neural Networks, Activation

```
tf.learn.Conv2D("Layer_1/Conv2D", Shape(3, 3, 3, 32), stride1 = 1, stride2 = 1, padding = ValidConvPadding) >>  
tf.learn.AddBias("Layer_1/Bias") >>  
tf.learn.ReLU("Layer_1/ReLU", alpha = 0.1f) >>
```


Rectifying Layer...Breaking positive/negative symmetry



Convolutional Neural Networks, subsampling

```
tf.learn.Conv2D("Layer_1/Conv2D", Shape(3, 3, 3, 32), stride1 = 1, stride2 = 1, padding = ValidConvPadding) >>  
tf.learn.AddBias("Layer_1/Bias") >>  
tf.learn.ReLU("Layer_1/ReLU", alpha = 0.1f) >>  
tf.learn.MaxPool("Layer_1/MaxPool", windowSize = Seq(1, 2, 2, 1), stride1 = 2, stride2 = 2) >>
```

MaxPooling: Scale down the 2-D size by selecting the best score per window pool

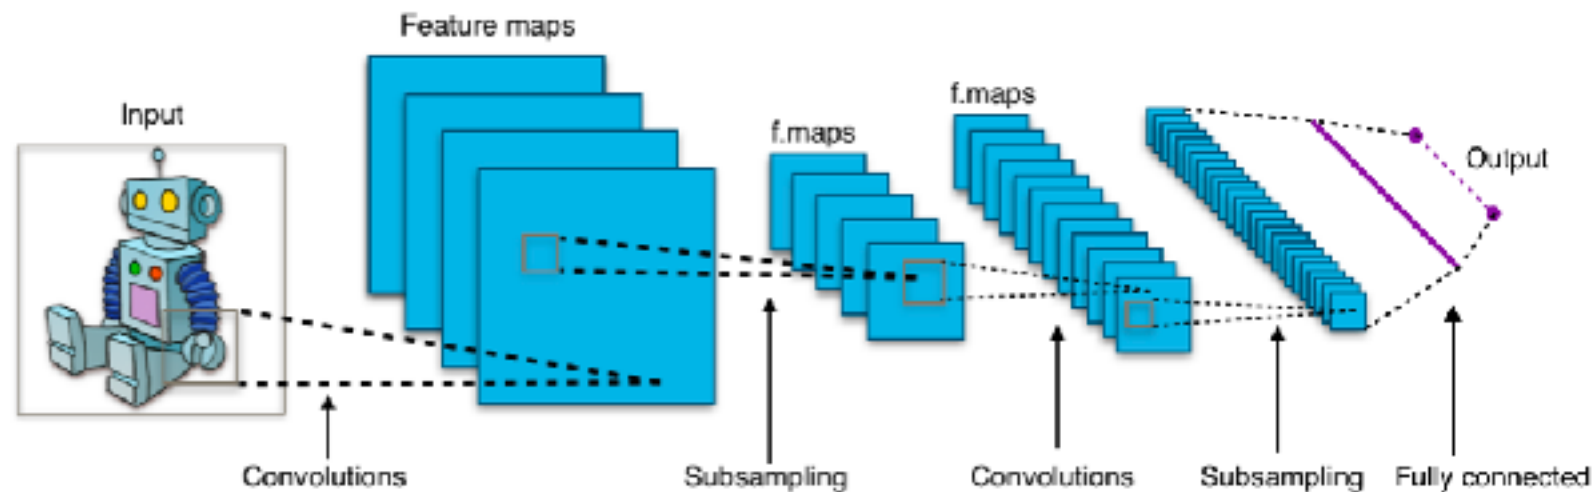


This creates one hierarchy level of detected patterns with locality

Convolutional Neural Networks

```
tf.learn.Conv2D("Layer_1/Conv2D", Shape(3, 3, 3, 32), stride1 = 1, stride2 = 1, padding = ValidConvPadding) >>
tf.learn.AddBias("Layer_1/Bias") >>
tf.learn.ReLU("Layer_1/ReLU", alpha = 0.1f) >>
tf.learn.MaxPool("Layer_1/MaxPool", windowSize = Seq(1, 2, 2, 1), stride1 = 2, stride2 = 2) >>
tf.learn.Flatten("Layer_2/Flatten") >>
tf.learn.Linear("Layer_2/Linear", units = 512) >>
tf.learn.ReLU("Layer_2/ReLU", 0.01f) >>
tf.learn.Linear("OutputLayer/Linear", 2)
```


And repeat...



Linear Models to Deep Learning

Convolutional Neural Network

```
tf.learn.Conv2D("Layer_1/Conv2D", Shape(3, 3, 3, 32), stride1 = 1, stride2 = 1, padding = ValidConvPadding) >>
tf.learn.AddBias("Layer_1/Bias") >>
tf.learn.ReLU("Layer_1/ReLU", alpha = 0.1f) >>
tf.learn.MaxPool("Layer_1/MaxPool", windowSize = Seq(1, 2, 2, 1), stride1 = 2, stride2 = 2) >>
tf.learn.Flatten("Layer_2/Flatten") >>
tf.learn.Linear("Layer_2/Linear", units = 512) >>
tf.learn.ReLU("Layer_2/ReLU", 0.01f) >>
tf.learn.Linear("OutputLayer/Linear", 2)
```

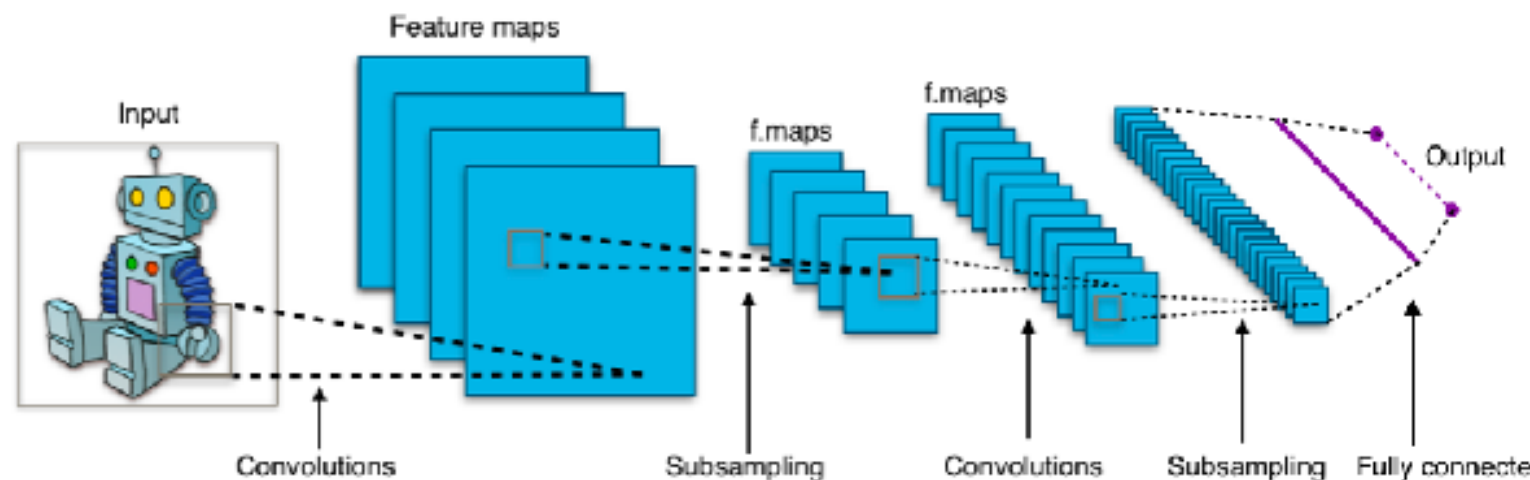


Deep is defined by accumulation of layers...and parameters

Linear Models to Deep Learning

Convolutional Neural Network

```
tf.learn.Conv2D("Layer_1/Conv2D", Shape(3, 3, 3, 32), stride1 = 1, stride2 = 1, padding = ValidConvPadding) >>  
tf.learn.AddBias("Layer_1/Bias") >>  
tf.learn.ReLU("Layer_1/ReLU", alpha = 0.1f) >>  
tf.learn.MaxPool("Layer_1/MaxPool", windowSize = Seq(1, 2, 2, 1), stride1 = 2, stride2 = 2) >>  
tf.learn.Flatten("Layer_2/Flatten") >>  
tf.learn.Linear("Layer_2/Linear", units = 512) >>  
tf.learn.ReLU("Layer_2/ReLU", 0.01f) >>  
tf.learn.Linear("OutputLayer/Linear", 2)
```



Deep is defined by accumulation of layers...and parameters

Tensorflow: More models

Resources for published models:

<https://github.com/tensorflow/models>

<https://tfhub.dev/>

Tensorflow: Takeaways

Data access API

Tensor representation and manipulation

Deeplearning Modelling API

Serializable Computing Graphs (Interoperable)

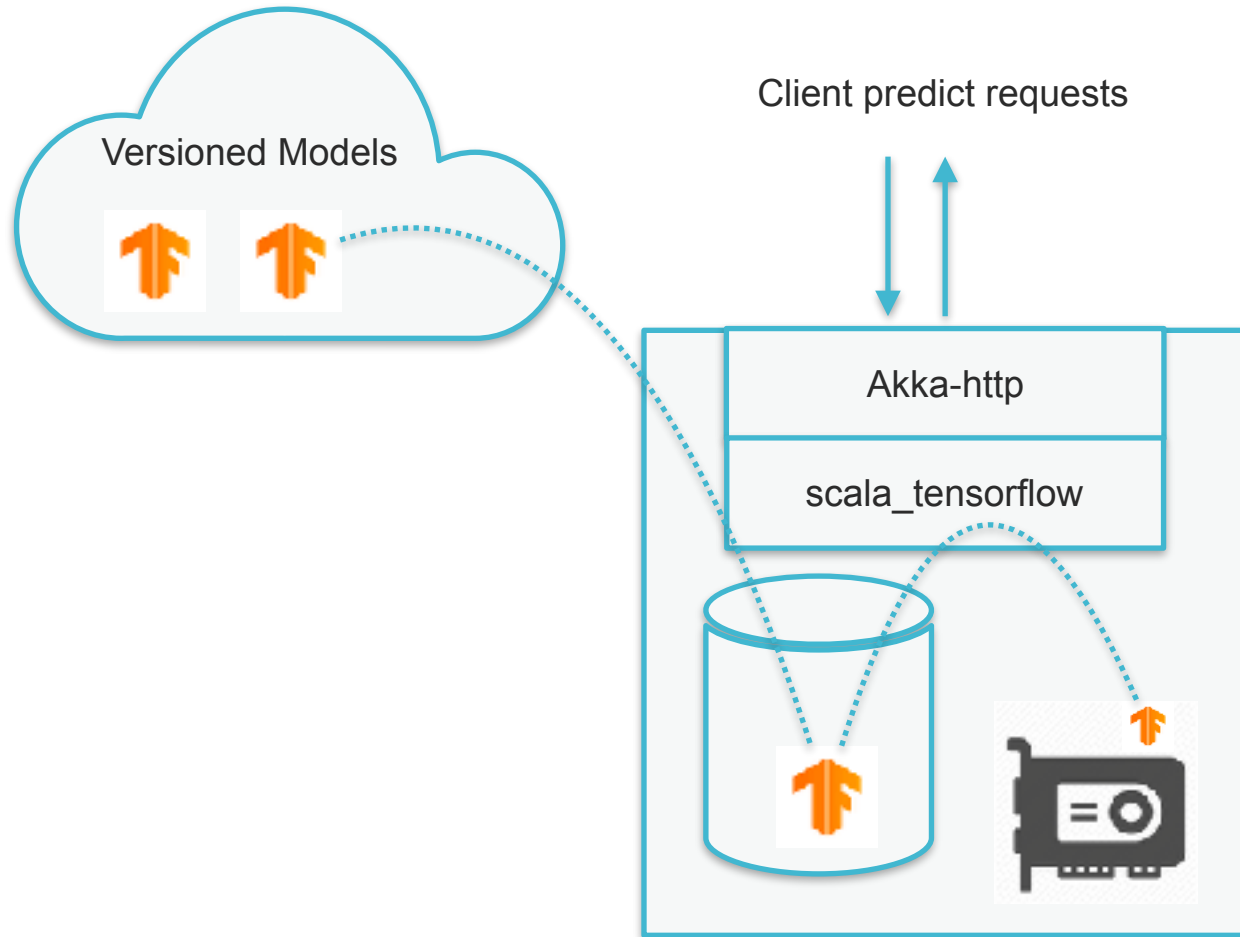
Access to Accelerated Computing Devices (gpu and tpu)

Tensorflow: Serving from Scala Demo

Goals

- Serve some common AI task with inherent complexity (e.g. object detection in images for anonymization)
- Leverage Scala strong concurrency and scheduling capabilities ... ***akka-http***
- Consume ML models in with a degree of abstraction: A single use case should have the same methods signatures and specific models easily interchangeable ... ***config-based model hub***

Tensorflow: Serving from Scala Demo



Goals

- **Serve** common **AI** task (e.g. object detection in images for anonymization)
- Leverage Scala concurrency and scheduling capabilities ... ***akka-http***
- Consume ML models: input signatures and interchangeable models ... ***config-based model hub***

Tensorflow: Serving from Scala Demo

Start in a notebook to build the app skeleton

02_akkahttp_tf.ipynb

03_tf_client.ipynb

Tensorflow: Serving from Scala Demo

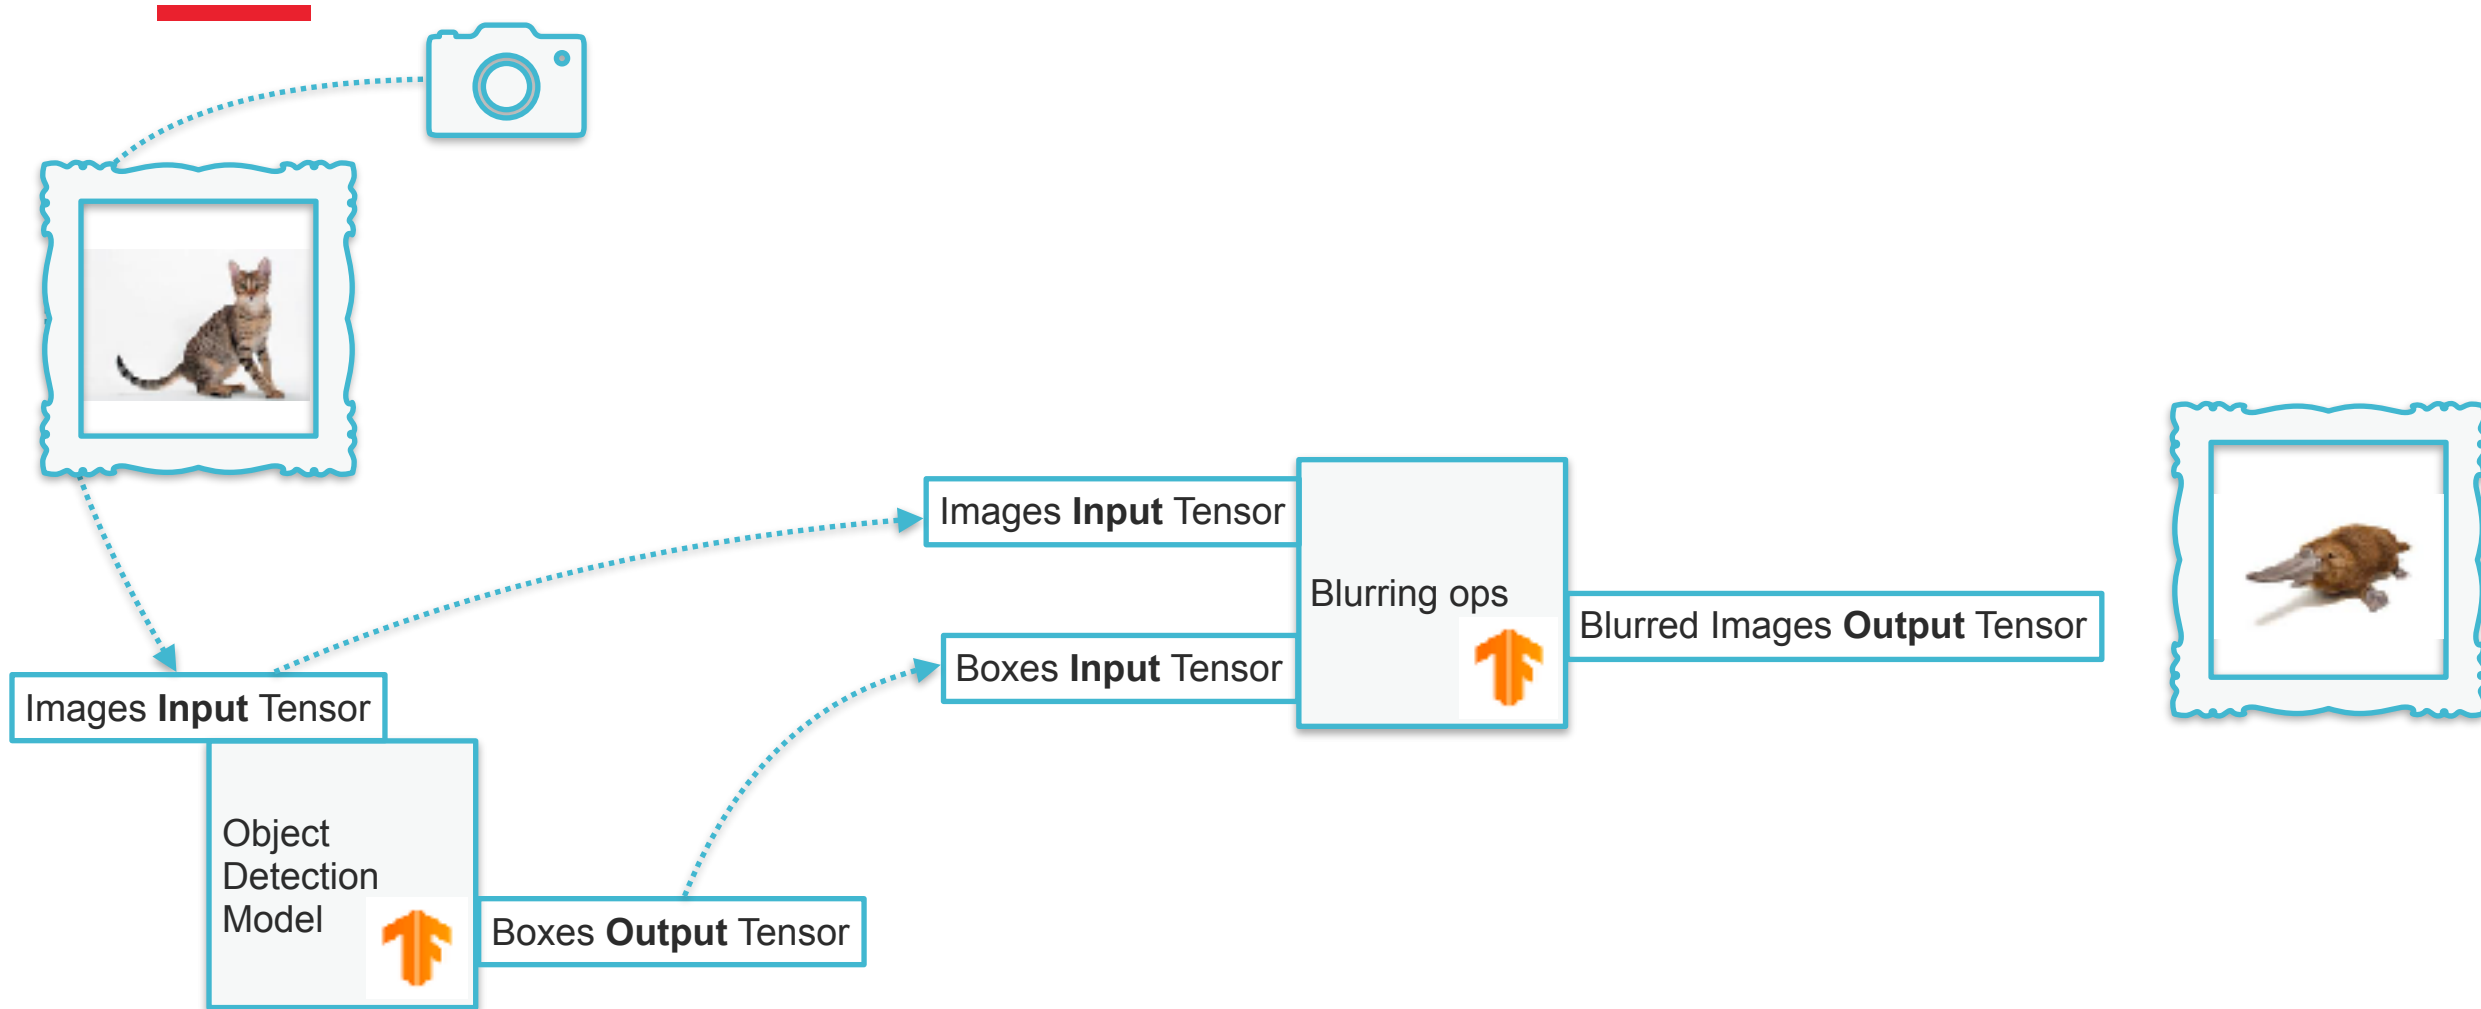
Move to sbt project

Abstract Model Service

Download from hub

Start service with different models

Tensorflow: Anonymization use case



Tensorflow: Unit Testing a Model

If the Tensorflow-side model signature is defined, how to unit test the code using it without relying on a large published model and some complex or unpredictable examples?

Just like database tables mocks, we need a mock of the model the can fit as test resources:

- Build a Tensorflow model with the required signature
- The Model should return predictable values for simple known cases
- And carry few computations

Example for object detection:

- return always the entire image and one defined quadrant
- class is always a pre-defined colour name
- score is the fraction of pixels with that exact colour

Such a model is small in size, and just requires a very small synthetic image to implement unit tests

Tensorflow: Integration Tests

Integration must use a real model from a hub, to check that Data Science team models are compatible with the requirements.

It would be bad that the signature changes even slightly with little transparency...

ML and Engineering: Team work

- Multiple environment and dependencies management is a burden to care about
- Favour Data Science artefacts as data, not code
- Define contracts: Model signatures, publishing protocols, Integration and Unit tests DS must comply with
- Leave room for experiments but value automation and maturity of models delivery

Scala Tensorflow alternatives

scalapy

<https://github.com/shadaj/scalapy>

JEP and python environment

Delicate threads management

Java API

Official Tensorflow API on JVM

Tensorflow serving

C++ compile / docker

Manages models versions

gRPC / REST

Actually...other engines (lite, js, python, rust, swift...)

General conclusions

ML Engineering is progressing

- Data is available
- Computing is available
- Software Engineering and Deployment is mature

Tensorflow brings important features:

- Serializable and interoperable Models
- Model signatures can be defined
- Serving gRPC ... the client doesn't even need tensorflow
- Tensorflow Hub defines policies to publish models in a consistent manner
- Tensorflow Extended provides full pipelines management

Scala doesn't need to be bitten by the snake when it comes to ML

Investing in getting lots done as Tensorflow operations is a valuable investment

Merci !

xavier.tordoir@lunatech.nl
twitter.com/xtordoir

www.lunatech.com



LUNATECH
SIMPLIFY YOUR IT